

Data Science ‘Mushrooms’ Capstone Project

András Gelencsér

14 December 2019

Overview

This document describes the result of the capstone project for the HarvardX Data Science course based on the Mushrooms dataset available in the UCI Machine Learning Repository (Dua, D. and Graff, C. (2019). UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science.)

The “Mushrooms” dataset was recorded down from The Audubon Society Field Guide to North American Mushrooms (1981). G. H. Lincoff (Pres.), New York: Alfred A. Knopf and is available on the <https://archive.ics.uci.edu/ml/datasets/mushroom> url.

The goal of this project is to develop a model based on data science and machine learning techniques to predict if a mushroom is eatable based on the known attributes.

The dataset contains 8124 records and uses 22 attributes to describe the mushrooms. Each dataset contains additional attribute to flag if the mushroom is eatable or not.

We will use 80% of the available data for train our model. The remaining 20% will be used as test dataset for evaluating the developed model.

For the developed model is very important to detect poisonous mushrooms to avoid possible health damages or death. Therefore we will evaluate the model performance with the F1 score instead of the overall accuracy, because it's very important to reach high sensitivity with our model.

The following steps were done during the project for developing the algorithm:

1. Download the data and prepare for the analysis
2. Analyze the data structure
3. Define the classification model based on the result the data analysis
4. Implement and train the model based on the train set
5. Review the model based on the test set

The final model performed an F1 score of 1 by using a prediction based on the 9th most important features.

Methods & analysis

Preparation

For the analysis the data we will use the following R packages:

- tidyverse
- caret
- gridExtra
- ggplot2

the following R code loads and install the required packages if needed:

```
if(!require(tidyverse)) install.packages("tidyverse",
                                          repos = "http://cran.us.r-project.org", dependencies = TRUE)
if(!require(caret)) install.packages("caret",
                                      repos = "http://cran.us.r-project.org", dependencies = TRUE)
if(!require(gridExtra)) install.packages("gridExtra",
```

```

      repos = "http://cran.us.r-project.org", dependencies = TRUE)
if(!require(ggplot2)) install.packages("ggplot2",
      repos = "http://cran.us.r-project.org", dependencies = TRUE)
if(!require(Rborist)) install.packages("Rborist",
      repos = "http://cran.us.r-project.org", dependencies = TRUE)

```

At first, we need to prepare the data set for the analysis. The mushrooms dataset contains 8,124 records about different mushrooms.

The following R code downloads the data from the <https://archive.ics.uci.edu/ml/datasets/mushroom> site. To avoid unnecessary data traffic, the data will be downloaded only if not done yet.

```

#create directory for the data file if necessary
if (!dir.exists("mushrooms")){
  dir.create("mushrooms")
}
baseurl = "https://archive.ics.uci.edu/ml/machine-learning-databases/mushroom/"
files_to_download = c("agaricus-lepiota.data", "agaricus-lepiota.names")

for (f in files_to_download){
  #download the files only if not done yet
  if (!file.exists(paste("./mushrooms/",f,sep=''))){
    {
      download.file(paste(baseurl,f,sep=''), paste("./mushrooms/",f,sep=''))
    }
  }
}

f1 <- file("mushrooms/agaricus-lepiota.data")
mushroom_data <- str_split_fixed(readLines(f1), ",", 23)
close(f1)

```

The dataset is stored in the “agaricus-lepiota.data” file. This file contains 23 columns separated via comma. There is no header line in the file. The description file “agaricus-lepiota.names” gives us information about the classes and the 22 attributes in the data file.

So we can extract the data by parsing the data file. The following R code reads the data and stores in a data frame:

```

f1 <- file("mushrooms/agaricus-lepiota.data")
mushroom_data <- str_split_fixed(readLines(f1), ",", 23)
close(f1)

#set the column names for the features
colnames(mushroom_data) <- c("classes", "cap-shape", "cap-surface", "cap-color",
                             "bruises?", "odor", "gill-attachment",
                             "gill-spacing", "gill-size", "gill-color", "stalk-shape",
                             "stalk-root", "stalk-surface-above-ring",
                             "stalk-surface-below-ring", "stalk-color-above-ring",
                             "stalk-color-below-ring", "veil-type", "veil-color",
                             "ring-number", "ring-type", "spore-print-color",
                             "population", "habitat")

#convert to data frame
mushroom_data <- as.data.frame(mushroom_data)

#remove temporary file variable
rm(f1)

```

For training the classification method we will use 80% of the available data as training set. The remaining data will be used for evaluating the developed method. The following R code creates the train and test set:

```
#initialize random sequenz
set.seed(1, sample.kind = "Rounding")
#create index for train and test set
#20% of the data will be used for the test set
test_idx = createDataPartition(y = mushroom_data$classes, times=1, p=0.2, list=FALSE)
train_data = mushroom_data[-test_idx,]
test_data = mushroom_data[test_idx,]
#remove temporary variables
rm(mushroom_data, test_idx)
```

Available Data

The dataset contains 8,124 records. For the analysis we have data about 6,500 mushrooms (80% of all data). All the features are categorical with different possible values.

1. cap-shape: bell=b,conical=c,convex=x,flat=f, knobbed=k,sunken=s
2. cap-surface: fibrous=f,grooves=g,scaly=y,smooth=s
3. cap-color: brown=n, buff=b, cinnamon=c, gray=g, green=r, pink=p, purple=u, red=e, white=w, yellow=y
4. bruises?: bruises=t, no=f
5. odor: almond=a, anise=l, creosote=c, fishy=y, foul=f, musty=m, none=n, pungent=p, spicy=s
6. gill-attachment: attached=a, descending=d, free=f, notched=n
7. gill-spacing: close=c, crowded=w, distant=d
8. gill-size: broad=b, narrow=n
9. gill-color: black=k, brown=n, buff=b, chocolate=h, gray=g, green=r, orange=o, pink=p, purple=u, red=e, white=w, yellow=y
10. stalk-shape: enlarging=e, tapering=t
11. stalk-root: bulbous=b, club=c, cup=u, equal=e, rhizomorphs=z, rooted=r, missing=?
12. stalk-surface-above-ring: fibrous=f, scaly=y, silky=k, smooth=s
13. stalk-surface-below-ring: fibrous=f, scaly=y, silky=k, smooth=s
14. stalk-color-above-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
15. stalk-color-below-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
16. veil-type: partial=p, universal=u
17. veil-color: brown=n, orange=o, white=w, yellow=y
18. ring-number: none=n, one=o, two=t
19. ring-type: cobwebby=c, evanescent=e, flaring=f, large=l, none=n, pendant=p, sheathing=s, zone=z
20. spore-print-color: black=k, brown=n, buff=b, chocolate=h, green=r, orange=o, purple=u, white=w, yellow=y
21. population: abundant=a, clustered=c, numerous=n, scattered=s, several=v, solitary=y
22. habitat: grasses=g, leaves=l, meadows=m, paths=p, urban=u, waste=w, woods=d

For testing the performance the prediction model we have about 1,600 mushrooms data (20% of all data) The dataset contains only the short 1 character code of the attribute values, but it's not disturbing for the model development.

Feature analysis

As first step we can investigate the distribution of the features corresponding to the edibility classification.

The following code plots this distribution for each feature.

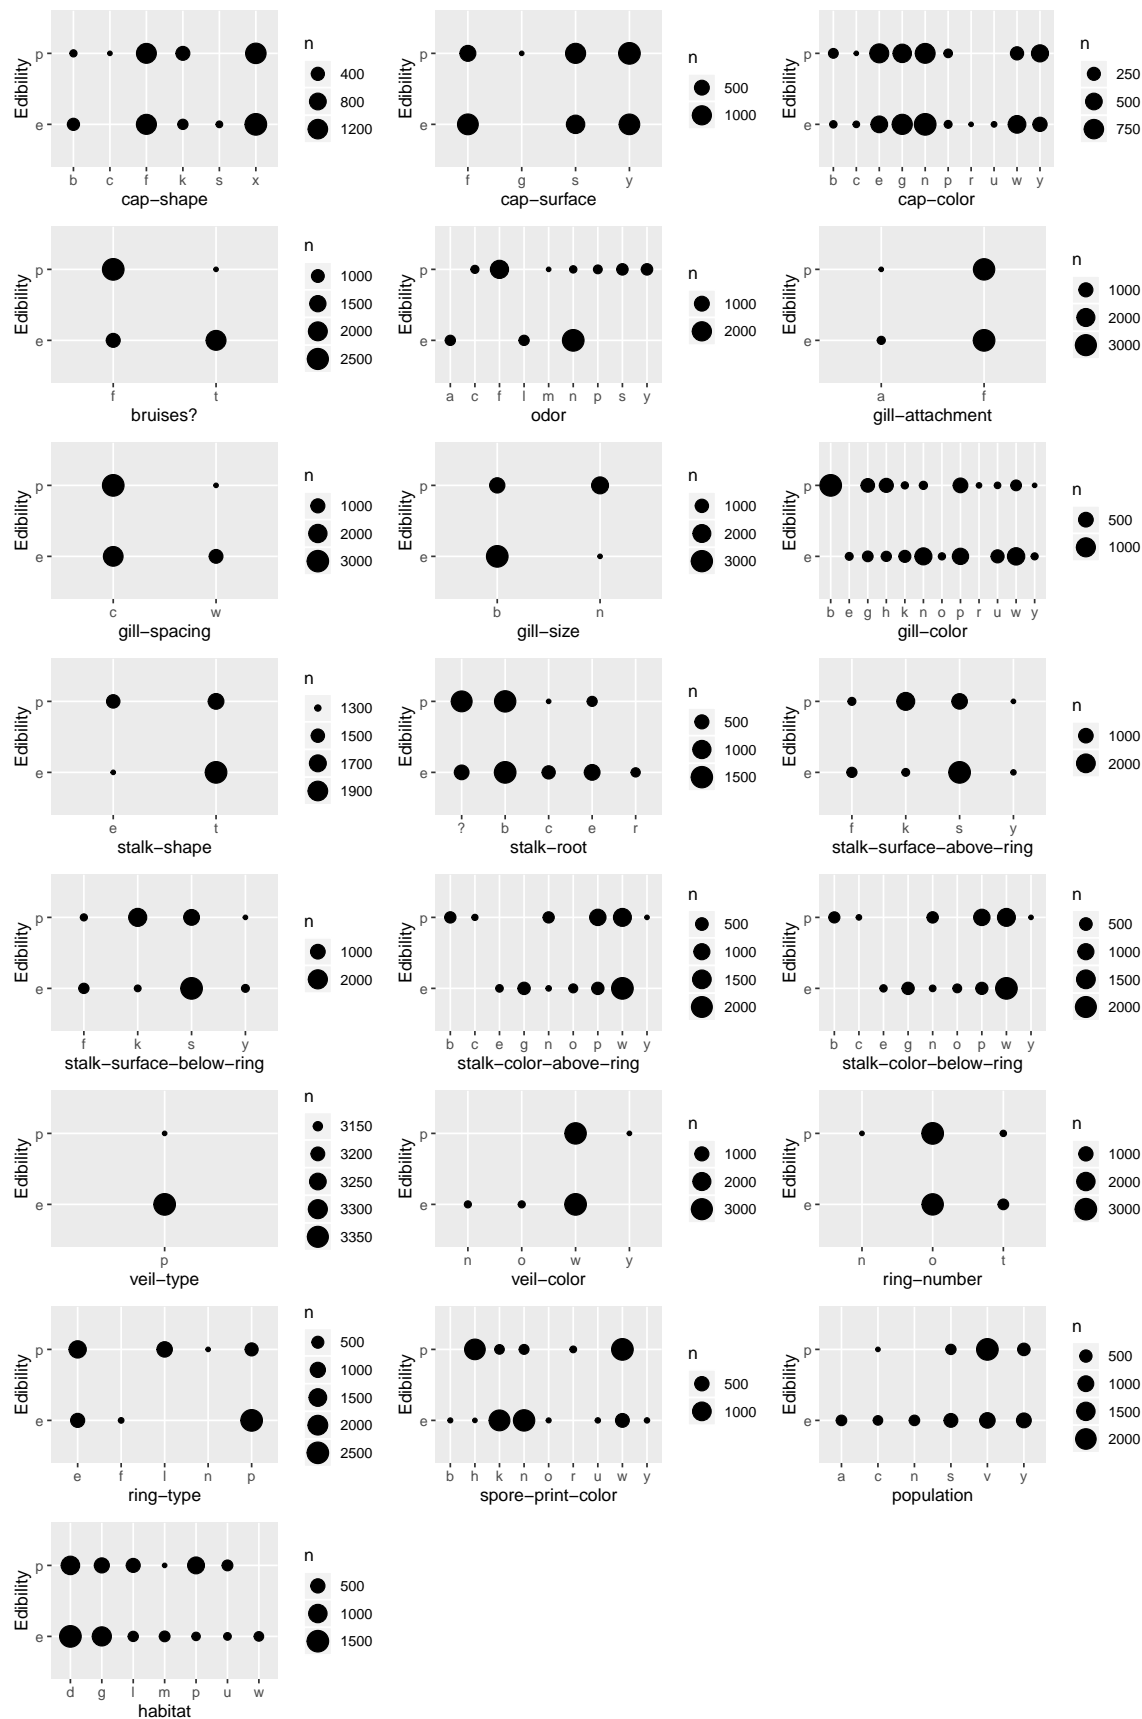
```
feature_distribution_plot <- function(data){
  plots <- list()
  #Plots all the attributes for eatable <-> poisonous
  for (i in 1:(ncol(data)-1))
```

```

{
  summarized_data <- data %>% group_by(classes, .[,i+1]) %>% summarise(n = n())
  names(summarized_data)[2] <- "attr"
  plot <- summarized_data %>% ggplot(aes(attr , classes)) + geom_point(aes(size=n)) +
    xlab(names(data)[i+1]) + ylab("Edibility")
  plots[[i]] <- plot
}
rm(summarized_data, i, plot)
grid.arrange(grobs=plots,ncol=3)
}

feature_distribution_plot(train_data)

```



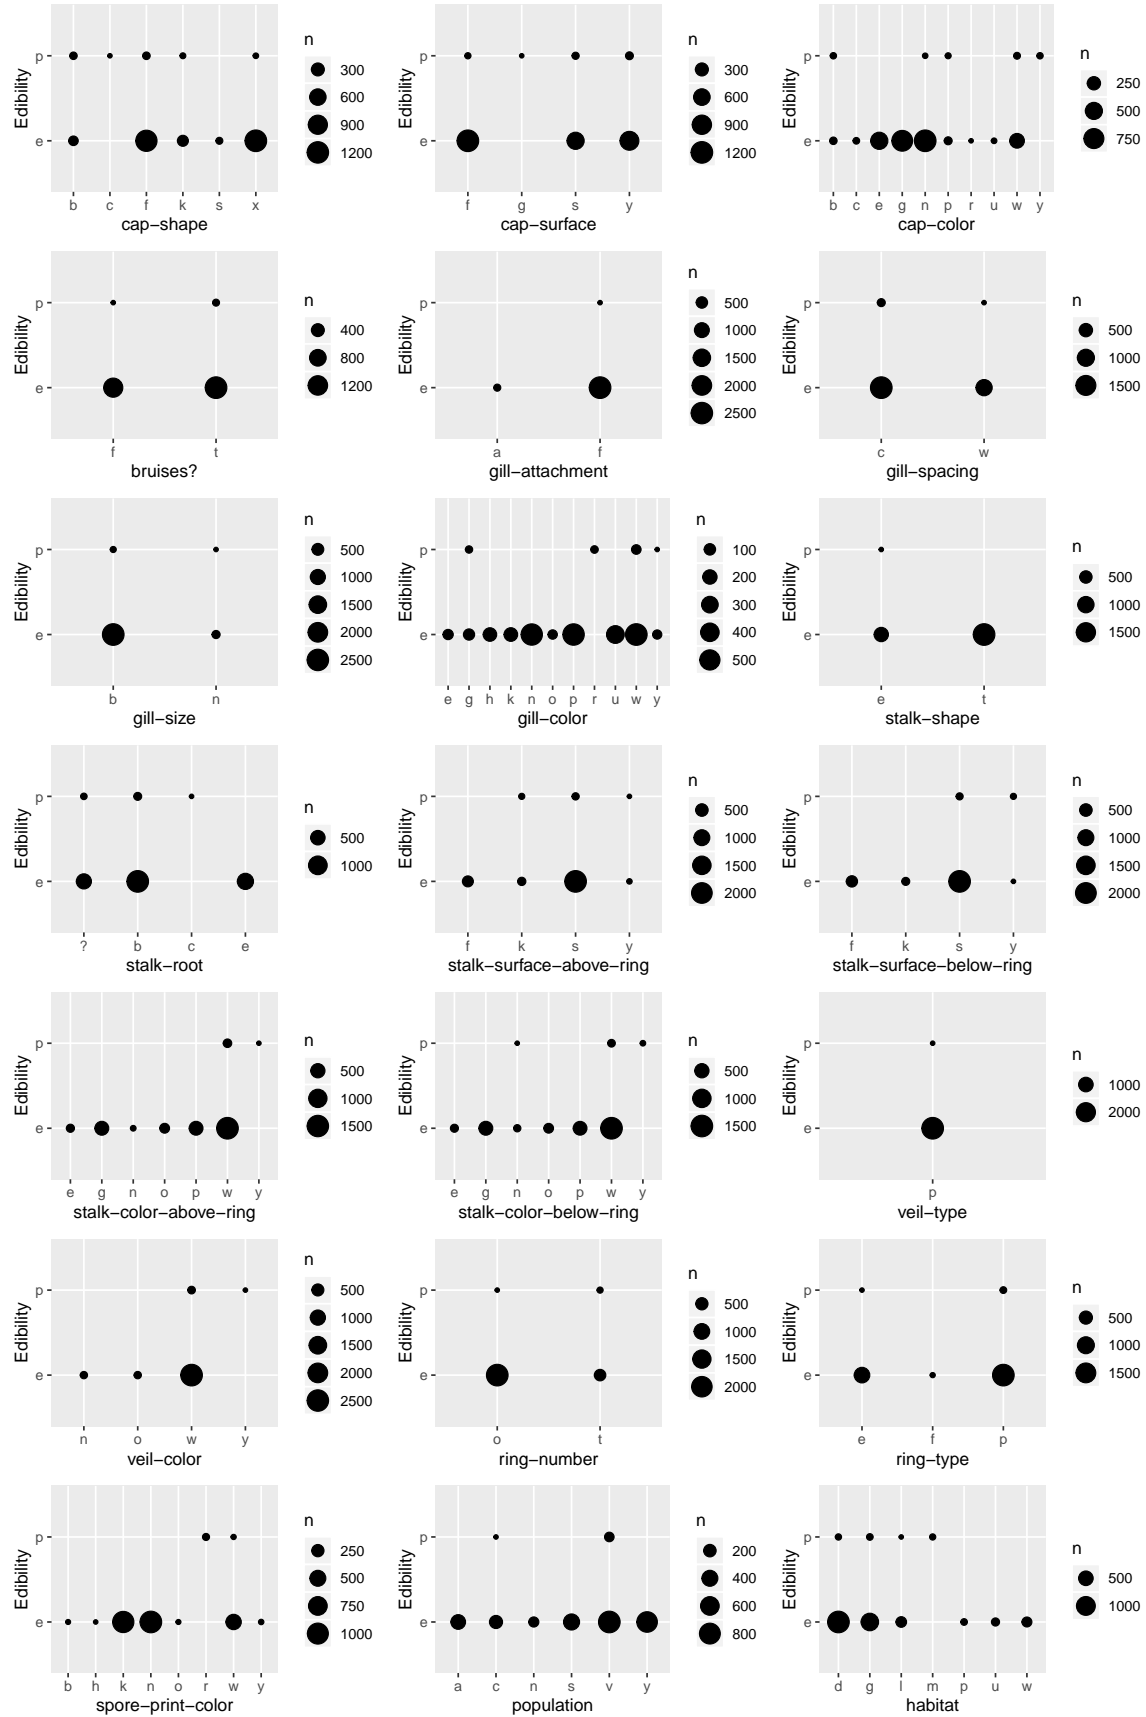
We can see, that the feature veil-type has only the value 'p' (partial), there is no data record with the veil-type value 'u' (universal). This feature can not provide decision information about the mushrooms, therefor we can remove this feature from the dataset.

On the other side, there are some features with values that can decide the classification. E.g. if we see a mushroom with the value 'f' (foul) for the feature odor, we can classify the mushroom as poisonous. The feature values 'c', 'f', 'l', 'm', 'p', 's', 'y' for the attribute odor can give us a classification. With this information we can classify 3,689 mushrooms. We can not classify only the mushrooms with the odor value 'n' (none). There are 2,809 mushrooms remaining in the dataset without classification.

We can try to analyse all the mushrooms, where the odor value is 'n'. For this analysis we can remove the attribute odor from the dataset, because we examine a specific odor value.

The following code removes the odor attribute and plots the feature distribution for the remaining data:

```
odor_n <- train_data %>% filter(`odor` == 'n') %>% select(-`odor`)
feature_distribution_plot(odor_n)
```



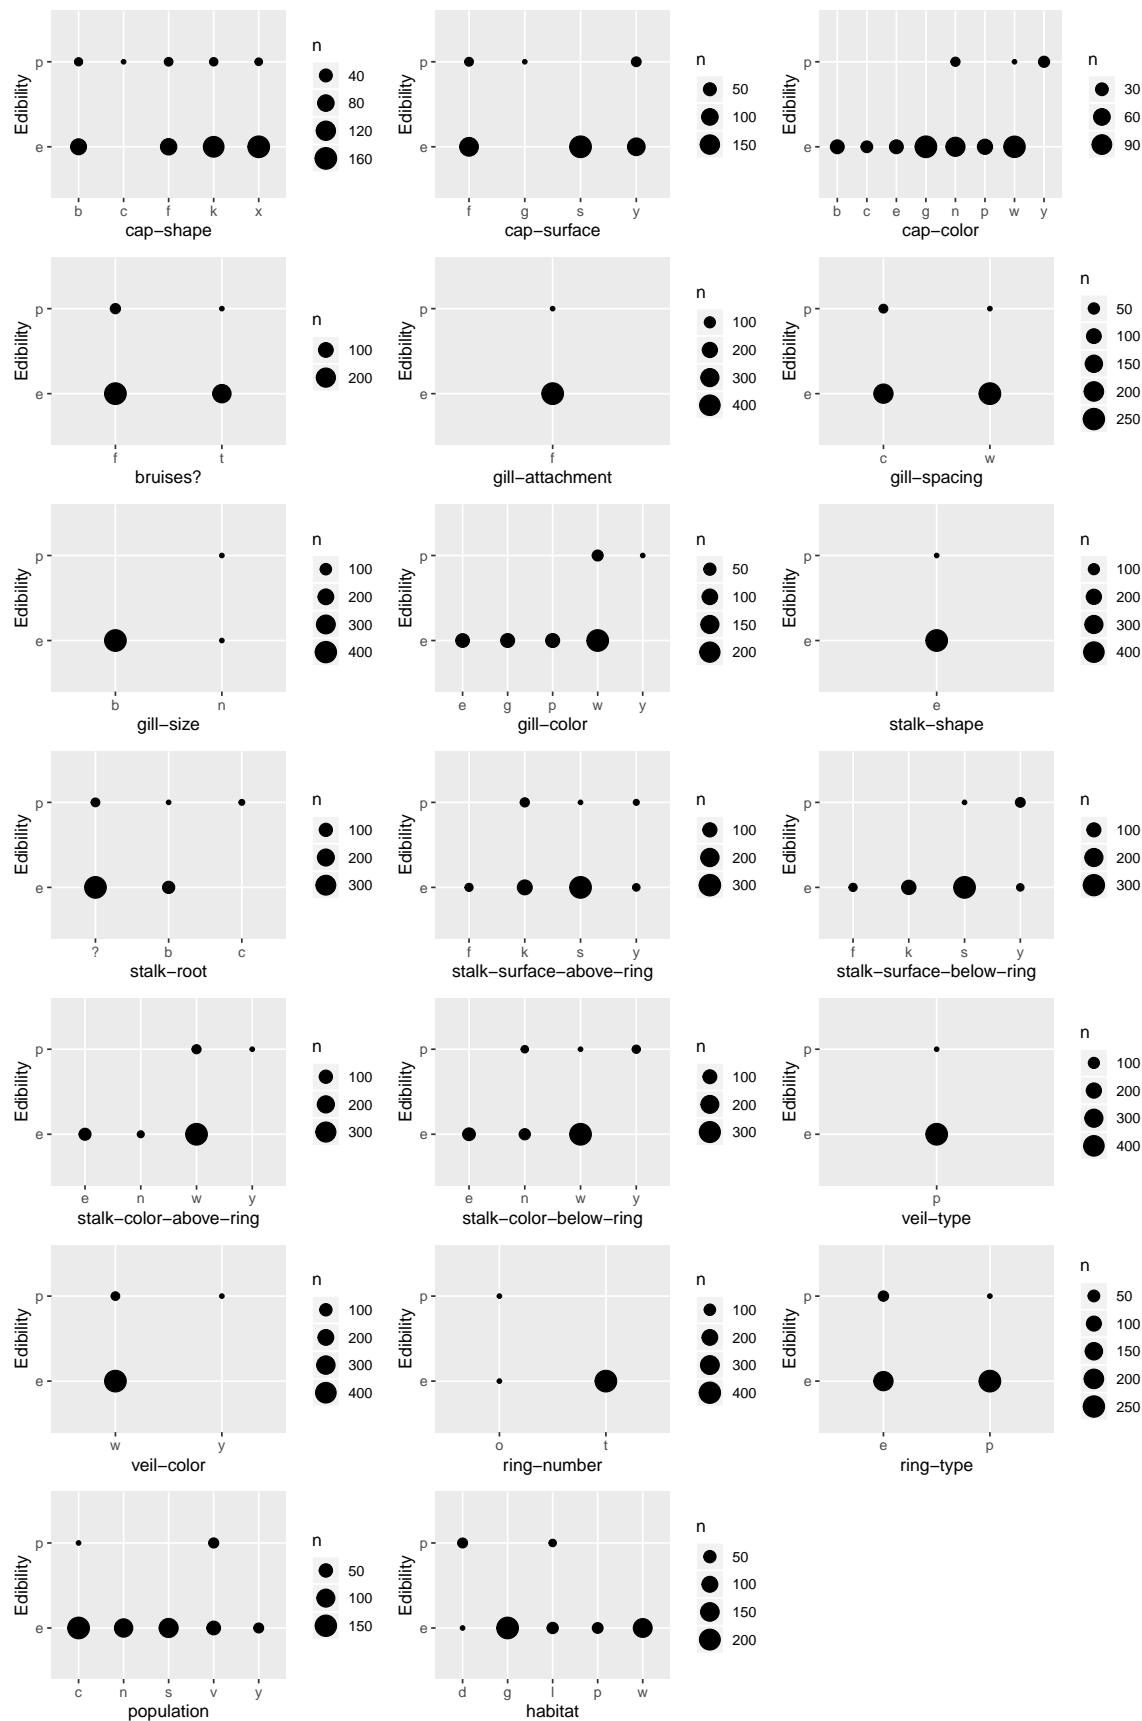
In the reduced data set we see again some features that we can use for classification the mushrooms.

E.g. if we see a mushroom with the value 'k' (black) for the feature spore-print-color, we can classify the mushroom as edible. The feature values 'b', 'h', 'k', 'n', 'o', 'r', 'y' for the feature spore-print-color can give us a classification for the remaining data set. We can not make a classification only if the spore-print-color value is 'w' (white).

We can reuse the same logic like for the feature odor and filter the remaining data with the spore-print-color feature value 'w'.

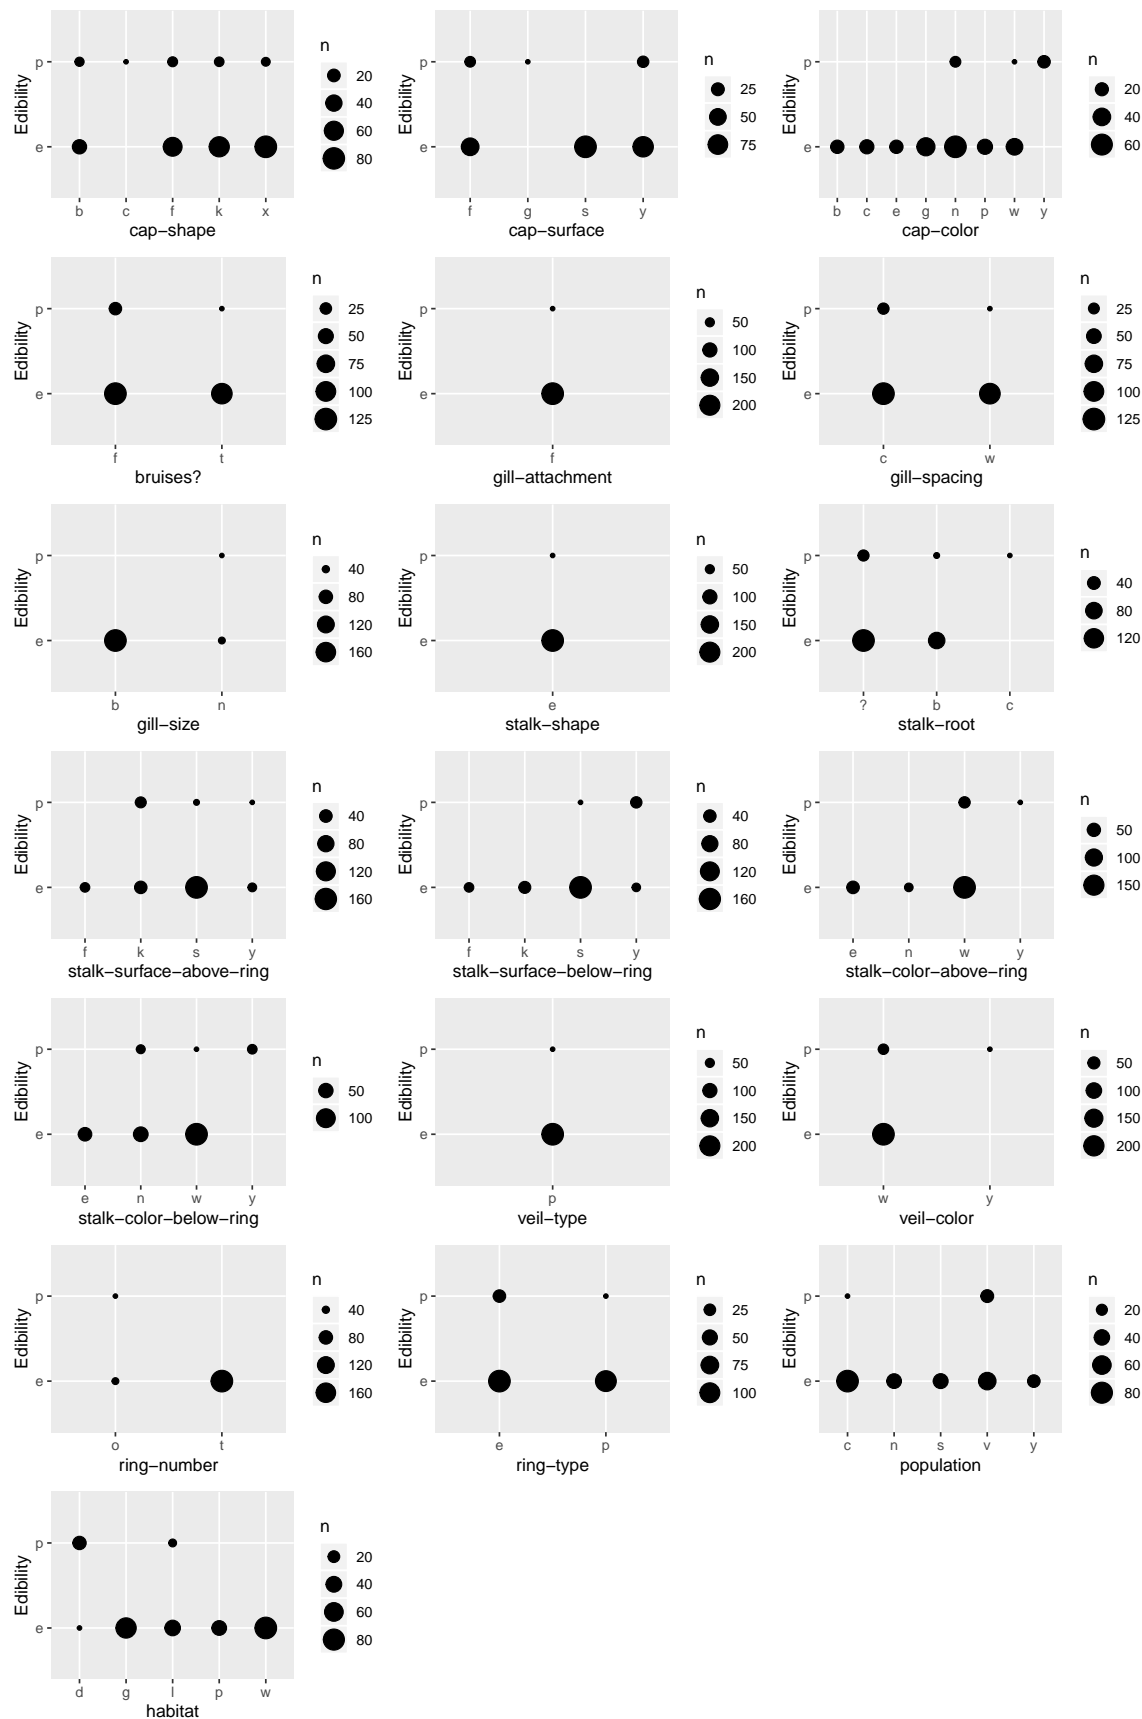
The following code removes the odor attribute and plots the feature distribution for the remaining data:

```
spore_print_color_w <- odor_n %>% filter(`spore-print-color` == 'w') %>%  
  select(-`spore-print-color`)  
feature_distribution_plot(spore_print_color_w)
```

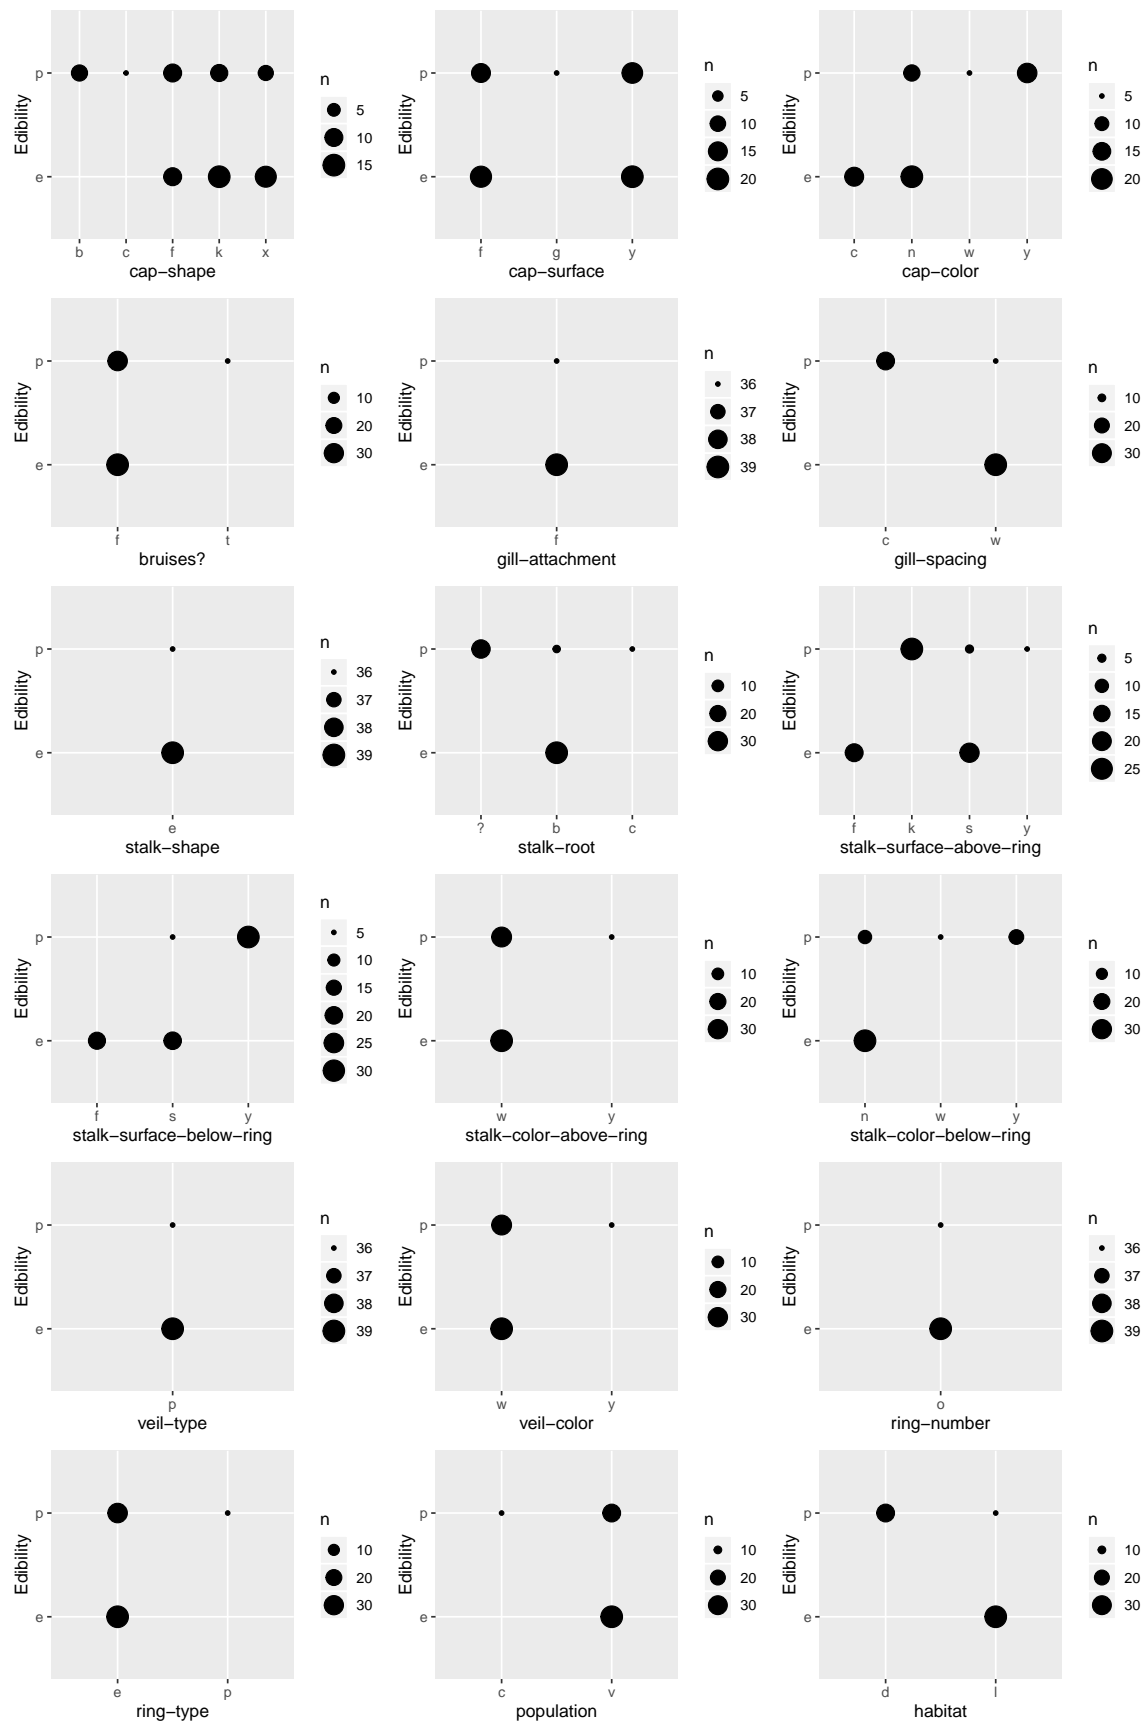
We can repeat this analysis and look for features that we can use for classifying the remaining data set. As next step we can get the gill-color feature and examine only the mushrooms with the gill color 'w' (white).

```
gill_color_w <- spore_print_color_w %>% filter(`gill-color` == 'w') %>%  
  select(-`gill-color`)  
feature_distribution_plot(gill_color_w)
```



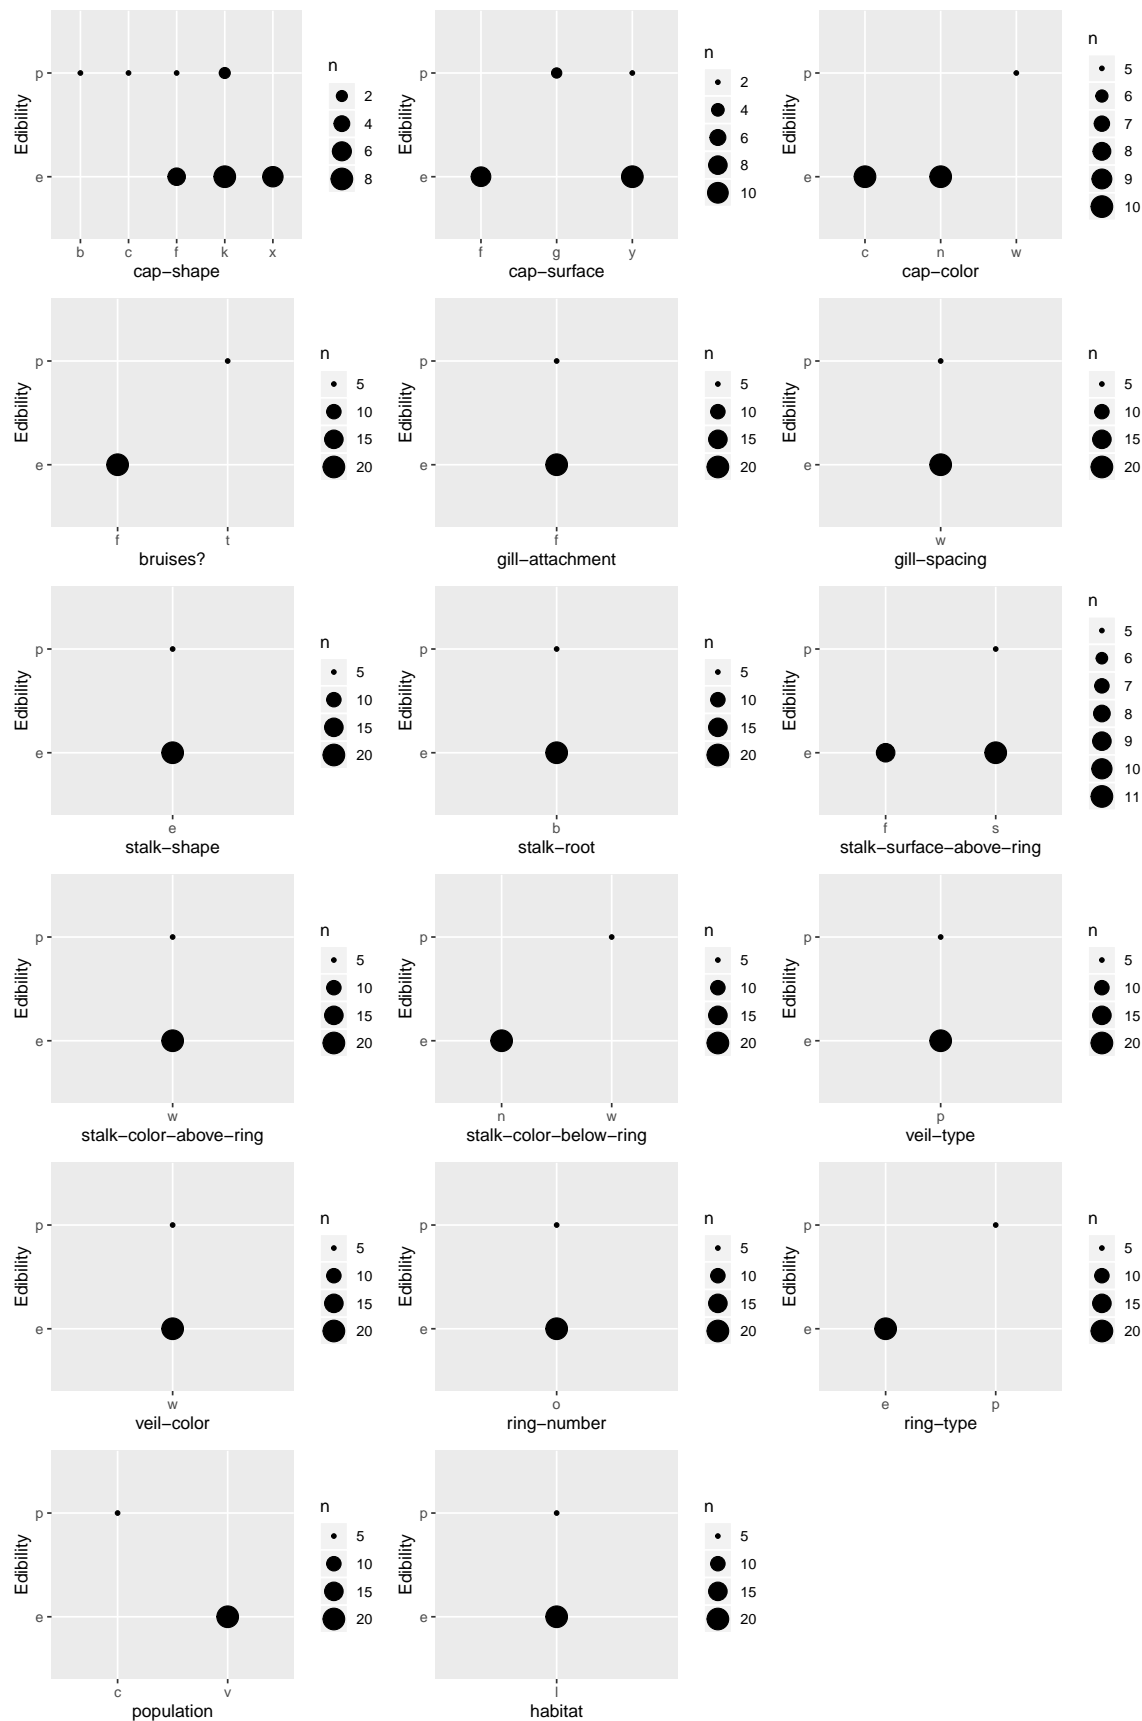
As next step we can get the gill-size feature and examine only the mushrooms with the grill size 'n' (narrow).

```
gill_size_n <- gill_color_w %>% filter(`gill-size` == 'n') %>% select(-`gill-size`)
feature_ditribution_plot(gill_size_n)
```



As next step we can get the stalk-surface-below-ring feature and examine only the mushrooms with the stalk-surface-below-ring 's' (smooth).

```
stalk_surface_below <- gill_size_n %>% filter(`stalk-surface-below-ring` == 's') %>%  
  select(-`stalk-surface-below-ring`)  
feature_distribution_plot(stalk_surface_below)
```



The re-

maining dataset has only 25 mushrooms. As we can see on the plot, now we can use some feature for classifying this remaining mushrooms. E.g. the feature ring-type give us a good classification: if we have a ring type of 'e' (evanescent) than the mushroom is edible. On the other case, if we have a ring-type of 'p' (pedant), the mushroom is poisonous.

With the steps about we defined a decision tree for classifying all the mushrooms in the train set. We can use this information to develop a classification model for the mushrooms.

The analyse above give us an other important information: some fatures are more relevant for the classification some features gives no additional information for the method. That means, the different features have different information weight for the classification.

We have to try to calculate a correlation between the features and the edibility. For the unordered categorical features we can use the uncertainty coefficient as correlation score. (See: https://en.wikipedia.org/wiki/Uncertainty_coefficient)

The uncertainty score builds on the conditinal entropy form the information theory.

The entropy for a single feature is defined as follows:

$$H(X) = - \sum_x P_X(x) \log P_X(x)$$

The conditional entropy of X given Y is defined as follows:

$$H(X|Y) = - \sum_{x,y} P_{X,Y}(x,y) \log P_{X|Y}(x|y)$$

The uncertainty coefficient of the two features can be calculated as follows:

$$U(X|Y) = \frac{H(X) - H(X|Y)}{H(X)}$$

We can implement an R function to calculate the entropy, conditional entropy and the uncertainty coefficient.

The following R code implements the calculation functions:

```
entropy <- function(dataset, colX){

  ret <- 0
  summarized_data <- dataset %>% group_by(.[,colX]) %>% summarise(n = n())
  rowCount <- nrow(dataset)
  level_items <- levels(dataset[,colX])
  for (item in level_items){
    prob <- summarized_data %>% filter(.[,1] == item) %>% pull(n) / rowCount
    ret <- ret + prob * log(prob,base = 2)
  }
  return (-ret)
}

cond_entropy <- function(dataset, colX, colY){
  ret <- 0
  rowCount <- nrow(dataset)
  summarized_y <- dataset %>% group_by(.[,colY]) %>% summarise(n = n())
  names(summarized_y)[1] <- "Y"
  level_itemsy <- levels(summarized_y$Y)

  for (itemy in level_itemsy){
    proby <- summarized_y %>% filter(Y == itemy) %>% pull(n) / rowCount
  }
}
```



```

summarized_x <- dataset %>% filter(.[,colY] == itemy) %>% group_by(.[,colX]) %>%
  summarise(n = n())
names(summarized_x)[1] <- "X"
level_itemsx <- levels(summarized_x$X)
for (itemx in level_itemsx){
  filtered <- summarized_x %>% filter(X == itemx)
  if (nrow(filtered) > 0 && probxy != 0){
    probxy <- filtered$n / rowCount
    probx_at_y <- probxy / proby
    ret <- ret + probxy * log(probx_at_y, base = 2)
  }
}
}
return (-ret)
}

uncertainty <- function(dataset, colX, colY){
  entr <- entropy(dataset,colX)
  cond_entr <- cond_entropy(dataset, colX, colY)
  if (entr == 0){
    return(0)
  }
  return ( (entr - cond_entr) / entr)
}

```

We can plot the uncertainty coefficient between all the features as a matrix. The following R code creates the plot:

```

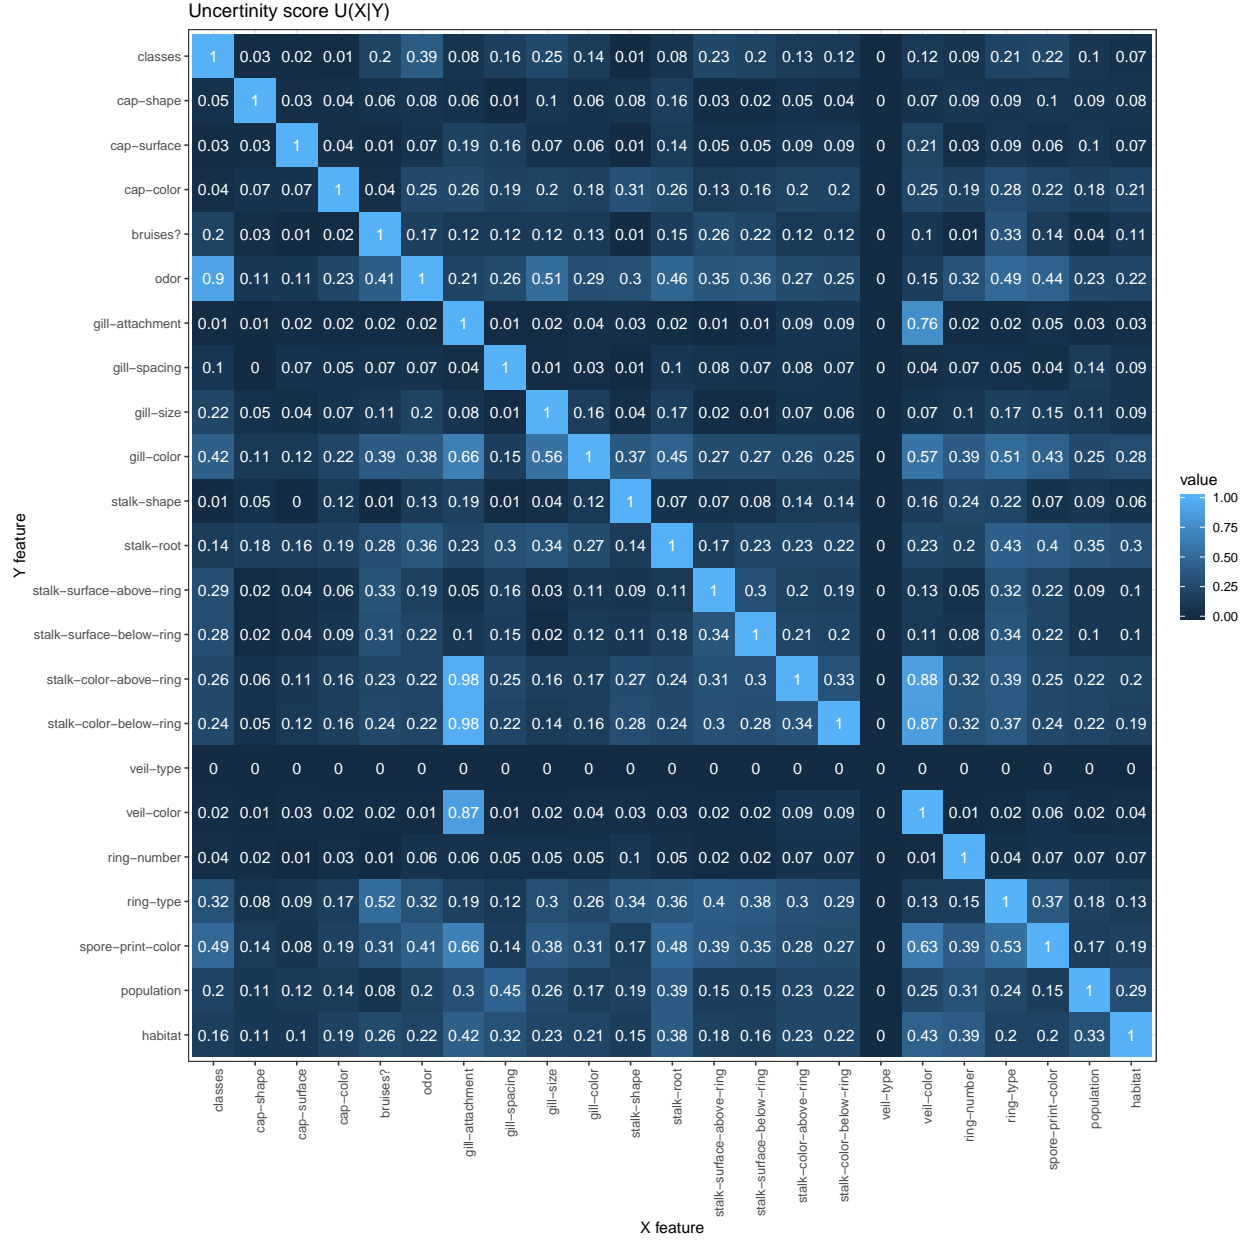
uncertainty_plot <- function(dataset){
  #calculate uncertainty for all feature pairs
  uncertainty_df <- data_frame(X=character(), idx=numeric(), Y=character(),
                              idy=numeric(), value=numeric())

  labels <- names(dataset)
  for (i in 1:ncol(dataset))
  {
    for(j in 1:ncol(dataset))
    {
      uncertainty_df <- bind_rows(uncertainty_df, data_frame(Y = labels[j], idy=j ,
                                                             X=labels[i], idx=i, value=uncertainty(train_data,i,j)))
    }
  }

  #plot the uncertainty with colors
  uncertainty_df %>% ggplot(aes(x=reorder(X,idx), y=reorder(Y,-idy), fill=value)) +
    geom_tile() + geom_text(aes(label=round(value,2)), color='white') +
    theme_bw() + theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
    xlab('X feature') + ylab('Y feature') + ggtitle('Uncertainty score U(X|Y)')
}

uncertainty_plot(train_data)

```



This plot shows the uncertainty coefficient between all feature pairs. The light blue values show a high uncertainty coefficient (high correlation) the dark blue values show a low uncertainty coefficient (low correlation). The value 1 in the matrix shows a perfect correlation, therefore the matrix diagonal elements are 1. (The diagonal elements show the correlation of the features with themselves, therefore there is a perfect correlation).

We see that the feature veil-type has 0 correlation with all other features. In the first data analysis we saw that the feature veil-type has only one value, that causes the 0 correlation.

For the classification model we examine the first column of the matrix. This column shows the correlation between edibility and other feature. As we see, the feature odor has the highest correlation with the edibility. The other features have only a correlation below 0.5.

We can use this information for developing a classification model that uses only the features with high correlation with the edibility.

The classification model

Decision tree model

Feature based model

Implementation & result

Conclusion